Springer Texts in Statistics

Gareth James
Daniela Witten
Trevor Hastie
Robert Tibshirani

An Introduction to Statistical Learning

with Applications in R

Second Edition

Springer

Gareth James · Daniela Witten ·
Trevor Hastie · Robert Tibshirani

An Introduction to Statistical Learning

with Applications in R

Second Edition

Springer

Contents

Preface vii

1 Introduction 1

2 Statistical Learning 15
- 2.1 What Is Statistical Learning? 15
  - 2.1.1 Why Estimate f? 17
  - 2.1.2 How Do We Estimate f? 21
  - 2.1.3 The Trade-Off Between Prediction Accuracy and Model Interpretability 24
  - 2.1.4 Supervised Versus Unsupervised Learning 26
  - 2.1.5 Regression Versus Classification Problems 28
- 2.2 Assessing Model Accuracy 29
  - 2.2.1 Measuring the Quality of Fit 29
  - 2.2.2 The Bias-Variance Trade-Off 33
  - 2.2.3 The Classification Setting 37
- 2.3 Lab: Introduction to R 42
  - 2.3.1 Basic Commands 43
  - 2.3.2 Graphics 45
  - 2.3.3 Indexing Data 47
  - 2.3.4 Loading Data 48
  - 2.3.5 Additional Graphical and Numerical Summaries . . 50
- 2.4 Exercises 52

3 Linear Regression 59
- 3.1 Simple Linear Regression 60
  - 3.1.1 Estimating the Coefficients 61
  - 3.1.2 Assessing the Accuracy of the Coefficient Estimates 63
  - 3.1.3 Assessing the Accuracy of the Model 68
- 3.2 Multiple Linear Regression 71
  - 3.2.1 Estimating the Regression Coefficients 72
  - 3.2.2 Some Important Questions 75
- 3.3 Other Considerations in the Regression Model 83

3.3.1 Qualitative Predictors 83
3.3.2 Extensions of the Linear Model 87
3.3.3 Potential Problems 92
3.4 The Marketing Plan 103
3.5 Comparison of Linear Regression with K-Nearest Neighbors 105
3.6 Lab: Linear Regression 110
3.6.1 Libraries 110
3.6.2 Simple Linear Regression 111
3.6.3 Multiple Linear Regression 114
3.6.4 Interaction Terms 116
3.6.5 Non-linear Transformations of the Predictors . . . 116
3.6.6 Qualitative Predictors 119
3.6.7 Writing Functions 120
3.7 Exercises 121

4 Classification **129**
4.1 An Overview of Classification 130
4.2 Why Not Linear Regression? 131
4.3 Logistic Regression 133
4.3.1 The Logistic Model 133
4.3.2 Estimating the Regression Coefficients 135
4.3.3 Making Predictions 136
4.3.4 Multiple Logistic Regression 137
4.3.5 Multinomial Logistic Regression 140
4.4 Generative Models for Classification 141
4.4.1 Linear Discriminant Analysis for $p = 1$ 142
4.4.2 Linear Discriminant Analysis for $p > 1$ 145
4.4.3 Quadratic Discriminant Analysis 152
4.4.4 Naive Bayes 153
4.5 A Comparison of Classification Methods 158
4.5.1 An Analytical Comparison 158
4.5.2 An Empirical Comparison 161
4.6 Generalized Linear Models 164
4.6.1 Linear Regression on the Bikeshare Data 164
4.6.2 Poisson Regression on the Bikeshare Data 167
4.6.3 Generalized Linear Models in Greater Generality . 170
4.7 Lab: Classification Methods 171
4.7.1 The Stock Market Data 171
4.7.2 Logistic Regression 172
4.7.3 Linear Discriminant Analysis 177
4.7.4 Quadratic Discriminant Analysis 179
4.7.5 Naive Bayes 180
4.7.6 K-Nearest Neighbors 181
4.7.7 Poisson Regression 185

4.8 Exercises 189

5 Resampling Methods 197
5.1 Cross-Validation 198
5.1.1 The Validation Set Approach 198
5.1.2 Leave-One-Out Cross-Validation 200
5.1.3 k-Fold Cross-Validation 203
5.1.4 Bias-Variance Trade-Off for k-Fold Cross-Validation 205
5.1.5 Cross-Validation on Classification Problems 206
5.2 The Bootstrap 209
5.3 Lab: Cross-Validation and the Bootstrap 212
5.3.1 The Validation Set Approach 213
5.3.2 Leave-One-Out Cross-Validation 214
5.3.3 k-Fold Cross-Validation 215
5.3.4 The Bootstrap 216
5.4 Exercises 219

6 Linear Model Selection and Regularization 225
6.1 Subset Selection 227
6.1.1 Best Subset Selection 227
6.1.2 Stepwise Selection 229
6.1.3 Choosing the Optimal Model 232
6.2 Shrinkage Methods 237
6.2.1 Ridge Regression 237
6.2.2 The Lasso 241
6.2.3 Selecting the Tuning Parameter 250
6.3 Dimension Reduction Methods 251
6.3.1 Principal Components Regression 252
6.3.2 Partial Least Squares 259
6.4 Considerations in High Dimensions 261
6.4.1 High-Dimensional Data 261
6.4.2 What Goes Wrong in High Dimensions? 262
6.4.3 Regression in High Dimensions 264
6.4.4 Interpreting Results in High Dimensions 266
6.5 Lab: Linear Models and Regularization Methods 267
6.5.1 Subset Selection Methods 267
6.5.2 Ridge Regression and the Lasso 274
6.5.3 PCR and PLS Regression 279
6.6 Exercises 282

7 Moving Beyond Linearity 289
7.1 Polynomial Regression 290
7.2 Step Functions 292
7.3 Basis Functions 294

7.4 Regression Splines 295
7.4.1 Piecewise Polynomials 295
7.4.2 Constraints and Splines 295
7.4.3 The Spline Basis Representation 297
7.4.4 Choosing the Number and Locations of the Knots 298
7.4.5 Comparison to Polynomial Regression 300
7.5 Smoothing Splines 301
7.5.1 An Overview of Smoothing Splines 301
7.5.2 Choosing the Smoothing Parameter λ 302
7.6 Local Regression 304
7.7 Generalized Additive Models 306
7.7.1 GAMs for Regression Problems 307
7.7.2 GAMs for Classification Problems 310
7.8 Lab: Non-linear Modeling 311
7.8.1 Polynomial Regression and Step Functions 312
7.8.2 Splines 317
7.8.3 GAMs 318
7.9 Exercises 321

8 Tree-Based Methods 327
8.1 The Basics of Decision Trees 327
8.1.1 Regression Trees 328
8.1.2 Classification Trees 335
8.1.3 Trees Versus Linear Models 338
8.1.4 Advantages and Disadvantages of Trees 339
8.2 Bagging, Random Forests, Boosting, and Bayesian Additive Regression Trees 340
8.2.1 Bagging 340
8.2.2 Random Forests 343
8.2.3 Boosting 345
8.2.4 Bayesian Additive Regression Trees 348
8.2.5 Summary of Tree Ensemble Methods 351
8.3 Lab: Decision Trees 353
8.3.1 Fitting Classification Trees 353
8.3.2 Fitting Regression Trees 356
8.3.3 Bagging and Random Forests 357
8.3.4 Boosting 359
8.3.5 Bayesian Additive Regression Trees 360
8.4 Exercises 361

9 Support Vector Machines 367
9.1 Maximal Margin Classifier 368
9.1.1 What Is a Hyperplane? 368
9.1.2 Classification Using a Separating Hyperplane . . . 369

9.1.3 The Maximal Margin Classifier 371
9.1.4 Construction of the Maximal Margin Classifier . . 372
9.1.5 The Non-separable Case 373
9.2 Support Vector Classifiers 373
9.2.1 Overview of the Support Vector Classifier 373
9.2.2 Details of the Support Vector Classifier 375
9.3 Support Vector Machines 379
9.3.1 Classification with Non-Linear Decision Boundaries 379
9.3.2 The Support Vector Machine 380
9.3.3 An Application to the Heart Disease Data 383
9.4 SVMs with More than Two Classes 385
9.4.1 One-Versus-One Classification 385
9.4.2 One-Versus-All Classification 385
9.5 Relationship to Logistic Regression 386
9.6 Lab: Support Vector Machines 388
9.6.1 Support Vector Classifier 389
9.6.2 Support Vector Machine 392
9.6.3 ROC Curves 394
9.6.4 SVM with Multiple Classes 396
9.6.5 Application to Gene Expression Data 396
9.7 Exercises 398

10 Deep Learning 403
10.1 Single Layer Neural Networks 404
10.2 Multilayer Neural Networks 407
10.3 Convolutional Neural Networks 411
10.3.1 Convolution Layers 412
10.3.2 Pooling Layers 415
10.3.3 Architecture of a Convolutional Neural Network . . 415
10.3.4 Data Augmentation 417
10.3.5 Results Using a Pretrained Classifier 417
10.4 Document Classification 419
10.5 Recurrent Neural Networks 421
10.5.1 Sequential Models for Document Classification . . 424
10.5.2 Time Series Forecasting 427
10.5.3 Summary of RNNs 431
10.6 When to Use Deep Learning 432
10.7 Fitting a Neural Network 434
10.7.1 Backpropagation 435
10.7.2 Regularization and Stochastic Gradient Descent . . 436
10.7.3 Dropout Learning 438
10.7.4 Network Tuning 438
10.8 Interpolation and Double Descent 439
10.9 Lab: Deep Learning 443

10.9.1 A Single Layer Network on the Hitters Data 443
10.9.2 A Multilayer Network on the MNIST Digit Data . 445
10.9.3 Convolutional Neural Networks 448
10.9.4 Using Pretrained CNN Models 451
10.9.5 IMDb Document Classification 452
10.9.6 Recurrent Neural Networks 454
10.10 Exercises 458

11 Survival Analysis and Censored Data 461
11.1 Survival and Censoring Times 462
11.2 A Closer Look at Censoring 463
11.3 The Kaplan-Meier Survival Curve 464
11.4 The Log-Rank Test 466
11.5 Regression Models With a Survival Response 469
11.5.1 The Hazard Function 469
11.5.2 Proportional Hazards 471
11.5.3 Example: Brain Cancer Data 475
11.5.4 Example: Publication Data 475
11.6 Shrinkage for the Cox Model 478
11.7 Additional Topics 480
11.7.1 Area Under the Curve for Survival Analysis 480
11.7.2 Choice of Time Scale 481
11.7.3 Time-Dependent Covariates 481
11.7.4 Checking the Proportional Hazards Assumption . . 482
11.7.5 Survival Trees 482
11.8 Lab: Survival Analysis 483
11.8.1 Brain Cancer Data 483
11.8.2 Publication Data 486
11.8.3 Call Center Data 487
11.9 Exercises 490

12 Unsupervised Learning 497
12.1 The Challenge of Unsupervised Learning 497
12.2 Principal Components Analysis 498
12.2.1 What Are Principal Components? 499
12.2.2 Another Interpretation of Principal Components . 503
12.2.3 The Proportion of Variance Explained 505
12.2.4 More on PCA 507
12.2.5 Other Uses for Principal Components 510
12.3 Missing Values and Matrix Completion 510
12.4 Clustering Methods 516
12.4.1 K-Means Clustering 517
12.4.2 Hierarchical Clustering 521
12.4.3 Practical Issues in Clustering 530
12.5 Lab: Unsupervised Learning 532

12.5.1 Principal Components Analysis 532
12.5.2 Matrix Completion 535
12.5.3 Clustering 538
12.5.4 NCI60 Data Example 542
12.6 Exercises 548

13 Multiple Testing **553**
13.1 A Quick Review of Hypothesis Testing 554
13.1.1 Testing a Hypothesis 555
13.1.2 Type I and Type II Errors 559
13.2 The Challenge of Multiple Testing 560
13.3 The Family-Wise Error Rate 561
13.3.1 What is the Family-Wise Error Rate? 562
13.3.2 Approaches to Control the Family-Wise Error Rate 564
13.3.3 Trade-Off Between the FWER and Power 570
13.4 The False Discovery Rate 571
13.4.1 Intuition for the False Discovery Rate 571
13.4.2 The Benjamini-Hochberg Procedure 573
13.5 A Re-Sampling Approach to p-Values and False Discovery Rates 575
13.5.1 A Re-Sampling Approach to the p-Value 576
13.5.2 A Re-Sampling Approach to the False Discovery Rate578
13.5.3 When Are Re-Sampling Approaches Useful? 581
13.6 Lab: Multiple Testing 582
13.6.1 Review of Hypothesis Tests 582
13.6.2 The Family-Wise Error Rate 583
13.6.3 The False Discovery Rate 586
13.6.4 A Re-Sampling Approach 588
13.7 Exercises 591

Index **597**

importance, 343, 358
independent, 15
indicator, 37
input, 15
output, 15
qualitative, 82–86, 89–91
selection, 78, 226, 242
variance, 19, 33–36, 155
inflation factor, 102–104, 115
varying coefficient model, 306
vector, 43

Wage data set, 1–3, 9, 10, 14, 291, 293, 295, 296, 298–301, 304, 306–308, 310, 311, 323, 324
weak learner, 340
weakest link pruning, 332
Weekly data set, 14, 193, 222
weight freezing, 419, 425
weight sharing, 423
weighted least squares, 97, 306
weights, 408
with replacement, 211
within class covariance, 146
workspace, 52
wrapper, 313